

# Package ‘talentlens’

July 6, 2016

**Type** Package  
**Title** R package for team TalentLens  
**Version** 0.2  
**Date** 2016-06-08  
**Author** Morgan Strom  
**Maintainer** <morgan.strom@pearson.com>  
**Description** Functions for test scoring, theta estimation etc.  
**License** No license  
**LazyData** TRUE  
**Depends** dplyr, tidyr, catR, psych, RMySQL, ltm

## R topics documented:

talentlens-package . . . . .	1
baggedParams . . . . .	2
balancer . . . . .	3
eap . . . . .	4
fetchConcertoData . . . . .	5
findAnchors . . . . .	6
monitorParticipants . . . . .	7
scoreDataFrame . . . . .	8
scoreResponses . . . . .	9
scoringKey . . . . .	10
targetUKCensus . . . . .	10
<b>Index</b>	<b>12</b>

---

talentlens-package	<i>Package for TalentLens R&amp;D team</i>
--------------------	--

---

## Description

Contains functions for getting access to data from a Concerto test, monitoring participants during data collection, balancing norm groups, finding optimal anchor items given a score matrix, scoring responses to items, estimating theta from a score matrix given item parameters, calculating bagged IRT parameters and generating target proportions according to UK census 2011.

Access information about these functions using the links below, or by typing e.g. `'?fetchConcertoData'` in the console.

## Details

Package: talentlens  
 Type: Package  
 Version: 0.2  
 Date: 2016-06-17  
 License: No licence

1. [fetchConcertoData](#)
2. [monitorParticipants](#)
3. [balancer](#)
4. [findAnchors](#)
5. [scoreDataFrame](#)
6. [scoreResponses](#)
7. [eap](#)
8. [baggedParams](#)
9. [targetUKCensus](#)

## Author(s)

Morgan Strom

---

baggedParams	<i>Function to calculate bagged IRT parameters from a score matrix</i>
--------------	--

---

## Description

Performs IRT fitting algorithm repeatedly (R times) on a random sample of the score matrix (with replacement), according to a bootstrap procedure. The results are then aggregated (bagging) and a list of parameter and error estimates is returned.

## Usage

```
baggedParams(score_mat, R, model = c("2pl", "rasch"),
  central.fun = mean, spread.fun = sd)
```

## Arguments

`score_mat` (nxk) numeric matrix rows corresponding to participants and columns corresponding to items. NOTE: columns need to be named after item ids.

R	Number of iterations for the bootstrap loop. Higher R increases the time the function needs for completion, but also gives better estimates of bagged parameters and bootstrapped errors. Default: 199 (on a Dell laptop, this takes about 1 min to run on a dataset with 99 rows).
model	The IRT model to use. This function currently supports rasch and 2pl models.
central.fun	The function used for aggregating results and find the central values. Default: mean.
spread.fun	The function used for estimating parameter errors. Default: sd.

**Value**

params	A (kx2) matrix of IRT parameters
error	A (kx2) matrix of estimation errors

**Author(s)**

Morgan Strom

---

balancer	<i>Function to balance a sample with regards to a dominant subgroup</i>
----------	---

---

**Description**

Takes a group variable (factor vector) *f* and the desired maximum proportion of the largest subgroup (factor level) *p*, samples randomly from the largest group and returns a logical vector, indicating which cases to retain to reach the desired proportion of the largest group.

**Usage**

```
balancer(f, p, silent = TRUE)
```

**Arguments**

f	Factor vector representing a group variable to be balanced, where the levels represent groups/categories in the data. For example, this could be a column in a dataframe, indicating the gender of the participants.
p	Desired maximum proportion for the largest subgroup. For example, if we want a balanced sample with regards to gender, we may want to allow for no more than 50 % of the total sample to be from either gender.
silent	If TRUE (default), messages are not printed to the console. Set to FALSE if you want to see how many cases are dropped randomly.

**Details**

Formula for determining the necessary reduction of the sample size (*delta\_n*) for the dominant subgroup:

$$\Delta n = \frac{N(p_2 - p_1)}{(1 - p_2)}$$

**Value**

Logical vector of the same length as `f`, with values `TRUE` for cases to keep and `FALSE` for cases to drop.

**Author(s)**

Morgan Strom

**Examples**

```
##Generate factor vector with unbalanced levels
f <- factor(x = c(rep("A", 20), rep("B", 50), rep("C", 30)))
table(f)

#Calculate proportions of the groups
table(f) / length(f)

#Use balancer() function to reduce the largest group to 40%

#First, create the index vector
i <- balancer(f, p = 0.40, silent = FALSE)

#Second, subset the factor vector and store in a new object, f2
f2 <- f[i]

#Calculate new proportions
table(f2) / length(f2)
```

---

eap

---

*Function to estimate ability using EAP algorithm*


---

**Description**

Takes a named score vector and a data frame with known item parameters, and returns the estimated theta.

**Usage**

```
eap(x, params, D = 1.702)
```

**Arguments**

<code>x</code>	Vector with correct responses coded as 1 and incorrect responses as 0. NOTE: the values need to be named with item ids
<code>params</code>	Data frame with unique item names and corresponding item parameters on each row, NOTE: columns need to be named "id", "a", "b" and "c"
<code>D</code>	Scaling constant D (Defaults to 1.702 for normal ogive model)

**Details**

The theta estimation algorithm is based on the implementation of the EAP algorithm in the Get-Feedback platform. Instructions from Louis-Charles Vannier.

An alternative to using this function is [eapEst](#) from the catR package.

**Value**

Estimated theta value.

**Author(s)**

Morgan Strom

**See Also**

[eapEst](#) [eapSem](#)

**Examples**

```
#Create named score vector
score_vec <- c(1,1,0,1,0)
names(score_vec) <- c("item1", "item2", "item3", "item4", "item5")

#Create parameter data frame
params <- data.frame(id = c("item1", "item2", "item3", "item4", "item5"),
  a = c(0.7, 0.8, 0.9, 1, 1.1),
  b = c(-2, -1, 0, 1, 2),
  c = c(0,0,0,0,0))

#Estimate theta
theta <- eap(score_vec, params)

#Using the function for a matrix with 4 observations
score_mat <- matrix(c(1,1,0,1,0,
  1,0,0,0,0,
  1,1,1,1,0,
  1,1,1,1,1),
  nrow = 4, ncol = 5, byrow = TRUE)
colnames(score_mat) <- c("item1", "item2", "item3", "item4", "item5")

thetas <- apply(score_mat, 1, eap, params=params)
```

---

fetchConcertoData

*Function to fetch data from Concerto*

---

**Description**

Returns candidate scores and demographics in a data frame format

**Usage**

```
fetchConcertoData(dbname, host, user, password, backup = FALSE, fetch = "scores")
```

**Arguments**

dbname	String containing the name of the database where the test is stored. E.g "concerto_2" for the UK version of Watson-Glaser. This information can be found in the spreadsheet "Concerto data collection" on Google Drive.
host	String containing the IP address of the Concerto server. This information can also be found in the spreadsheet "Concerto data collection" on Google Drive.
user	String containing the user name for accessing the database. To get access to this information, please email the author of this package.
password	String containing the password for accessing the database. To get access to this information, please email the author of this package.
backup	Boolean value - if TRUE (default), all candidate information from the server (tables candidate_summary and candidate_responses) will be stored as .csv files in the working directory.
fetch	String defining what data to fetch. If "scores" (default), the data will contain the scores matrix with 1 representing a correct response and 0 representing an incorrect response. If "responses", the data will instead contain participants' responses to the items.

**Value**

complete_data	A dataframe containing both demographic information and item scores from the concerto test. If you stored this in a dataframe called df, and you wish to save this as a .csv file in your working directory, type <code>write.csv(df, file = "df.csv", na = "", row.names = NULL)</code> in the console and press Enter.
---------------	--

**Author(s)**

Morgan Strom, TalentLens UK

---

findAnchors

---

*Function to suggest optimal anchor items from a score matrix*


---

**Description**

This function selects optimal anchor items in two steps: first, based on CTT analyses (proportion of participants answering the item correctly and item-total biserial correlation) and second, based on a chi-square test of item fit under the chosen IRT model.

**Usage**

```
findAnchors(score_matrix, model = c("rasch", "2pl"),
  bagging = FALSE, R = NA,
  min_cor = 0.2, min_p = 0.1, max_p = 0.9, verbose = FALSE)
```

**Arguments**

score_matrix	A numeric matrix where each row represent the scores of a participant, and each row represent a item. NOTE the columns need to be named with item ids!
model	The IRT model to use for parameter fitting. Either "rasch" (for smaller sample sizes) or "2pl" (for larger sample sizes)
bagging	If TRUE, the parameters will be estimated using bagging (see package bagged-Params). If FALSE, parameters will be estimated using one single run of the item fitting algorithm.
R	For bagging, this is the number of bootstrap iterations.
min_cor	The minimum allowed item-total biserial correlation. Default: 0.2
min_p	The minimum allowed proportion of participants answering a given item correctly. Default: 0.1
max_p	The maximum allowed proportion of participants answering a given item correctly. Default: 0.9
verbose	If TRUE, information about the item selection will be printed to the console.

**Value**

A two-column matrix of item parameters for the suggested anchor items. The rows are named after the item ids, and the columns are named "b" (difficulty parameter) and "a" (discrimination parameter).

**Author(s)**

Morgan Strom

---

monitorParticipants	<i>Function to monitor participants from MTurk projects</i>
---------------------	---

---

**Description**

Provides MTurk Worker ID along with time of completion, number of attempted items and total time for completion (in minutes).

**Usage**

```
monitorParticipants(dbname, host, user, password)
```

**Arguments**

dbname	String containing the name of the database where the test is stored. E.g "concerto_2" for the UK version of Watson-Glaser. This information can be found in the spreadsheet "Concerto data collection" on Google Drive.
host	String containing the IP address of the Concerto server. This information can also be found in the spreadsheet "Concerto data collection" on Google Drive.
user	String containing the user name for accessing the database. To get access to this information, please email the author of this package.
password	String containing the password for accessing the database. To get access to this information, please email the author of this package.

**Value**

Data frame containing the columns session\_id, worker\_id, timestamp, attempted and total\_time.

**Author(s)**

Morgan Strom

---

scoreDataFrame	<i>Function to score an entire data frame</i>
----------------	---

---

**Description**

Loops over all items in "key" and returns a score matrix with named columns. NOTE: the rows will be in the same order as the original data frame "df", but the columns may be in a different order.

**Usage**

```
scoreDataFrame(df, key)
```

**Arguments**

df	A data frame containing one candidate per row and one item per column (may also include additional columns, like candidate id etc). NOTE: the column names for the items need to correspond to the item ids in the column "id" in the "key" data frame.
key	A data frame containing one column for item ids ("id") and one column for the correct responses ("key"). If the correct response is a range, this should be represented in the following format: "RANGE: 1 - 1.1".

**Value**

A numeric score matrix, where the rows correspond to participants (same order as in the original data frame df) and the columns correspond to item ids (may be in different order from the original data frame). 1 represents a correct response, 0 an incorrect response and NA represents missing data.

**Author(s)**

Morgan Strom

**Examples**

```
key <- data.frame(id = c("item_1", "item_2", "item_3"),
                 key = c("2", 3, "RANGE: 1 - 1.1"))

df <- data.frame(id = c("cand_1", "cand_2", "cand_3"),
                item_1 = c(2, 1.99, 2.01),
                item_2 = c(3, 2.99, 3.01),
                item_3 = c(1, 0.99, 1.01))

scoreDataFrame(df, key)
```



---

scoreResponses	<i>Function to score a vector of responses</i>
----------------	--

---

### Description

Returns a vector of scores, given a vector of responses to a specific item, a vector with the correct response(s) ("key") and the item id.

### Usage

```
scoreResponses(responses, key, item_id)
```

### Arguments

responses	A numeric vector with the responses to an item
key	A numeric vector with the correct response to the item. If the response format is Multiple Choice, this will be a single integer value. If the response format is Free (such as some items in Athena), this will be two values representing the lower and the upper bounds of the range of acceptable responses.
item_id	A character vector with the item id

### Value

A numeric score vector with 1 for correct responses, 0 for incorrect responses and NA for missing values

### Author(s)

Morgan Strom

### Examples

```
responses <- c(1, 0, 1.05, 0.99, 2)

#Multiple choice with 2 as the correct answer
key <- 2
scoreResponses(responses, key, "test_item")

#Free response with lower bound 1 and upper bound 1.1
key <- c(1, 1.1)
scoreResponses(responses, key, "test_item")
```

---

scoringKey	<i>Function to create a scoring key for the scoreData function</i>
------------	--

---

### Description

Creates a list with one entry per item, where the name is the item id and the value is the correct response. The value can either be a single number, when there is only one correct answer, or a vector of two numbers, where the correct answer is within a range (Athena items for example)

### Usage

```
scoringKey(key.df)
```

### Arguments

key.df	Data frame containing one column with item ID (named "id") and one column with the corresponding correct responses (named "key"). If the correct response is a range, use the following format: "RANGE: 1 - 1.1".
--------	---

### Value

List with one entry per item. The entry is named after the item id, and the value is the correct response to the item.

### Author(s)

Morgan Strom

### Examples

```
df <- data.frame(id = c("item_1", "item_2", "item_3"),
  key = c("2", 3, "RANGE: 1 - 1.1"))

scoringKey(df)
```

---

targetUKCensus	<i>Function to calculate target proportions from UK Census data 2011</i>
----------------	--

---

### Description

Returns a dataframe with 60 rows, with target proportions for each possible subgroup with regards to age, gender, ethnicity and education level

### Usage

```
targetUKCensus()
```

**Details**

Age is divided into 2 groups: 16-24 and 25-64. Gender is also divided into 2 groups: Female and Male. Ethnicity is divided into 5 groups: White, Mixed, Asian, Black/African/Caribbean (BAC) and Other. Finally, Education is divided into 3 groups: No University Education (Levels 1-3), University Education (Level 4) and Other.

**Value**

Data frame containing all possible combinations of the groups described above ( $2 * 2 * 5 * 3 = 60$  rows), where each row represent a unique combination. The column "ExpectedP" contains the target proportion for that specific intersection of groups.

**Author(s)**

Morgan Strom

# Index

- \*Topic **\textasciitildeIO**
  - fetchConcertoData, [5](#)
  - monitorParticipants, [7](#)
- \*Topic **\textasciitildekwd1**
  - baggedParams, [2](#)
- \*Topic **\textasciitildemisc**
  - balancer, [3](#)
  - eap, [4](#)
  - fetchConcertoData, [5](#)
  - findAnchors, [6](#)
  - monitorParticipants, [7](#)
  - scoreDataFrame, [8](#)
  - scoreResponses, [9](#)
  - scoringKey, [10](#)
  - targetUKCensus, [10](#)
- \*Topic **package**
  - talentlens-package, [1](#)

baggedParams, [2](#), [2](#)  
balancer, [2](#), [3](#)

eap, [2](#), [4](#)  
eapEst, [5](#)  
eapSem, [5](#)

fetchConcertoData, [2](#), [5](#)  
findAnchors, [2](#), [6](#)

monitorParticipants, [2](#), [7](#)

scoreDataFrame, [2](#), [8](#)  
scoreResponses, [2](#), [9](#)  
scoringKey, [10](#)

talentlens (talentlens-package), [1](#)  
talentlens-package, [1](#)  
targetUKCensus, [2](#), [10](#)